

```
In [ ]: """  
Data Analytics II  
  
1. Implement logistic regression using Python/R to perform classification on  
Social_Network_Ads.csv dataset.  
  
2. Compute Confusion matrix to find TP, FP, TN, FN, Accuracy, Error rate, Precision, Recall  
on the given dataset.  
"""
```

```
In [10]: # Importing the dependencies  
  
import pandas as pd  
import numpy as np  
import seaborn as sns  
import matplotlib.pyplot as plt
```

```
In [2]: # Importing dataset from local file  
  
df = pd.read_csv("Social_Network_Ads.csv")
```

In [3]:

df

Out[3]:

	User ID	Gender	Age	EstimatedSalary	Purchased
0	15624510	Male	19	19000	0
1	15810944	Male	35	20000	0
2	15668575	Female	26	43000	0
3	15603246	Female	27	57000	0
4	15804002	Male	19	76000	0
...
395	15691863	Female	46	41000	1
396	15706071	Male	51	23000	1
397	15654296	Female	50	20000	1
398	15755018	Male	36	33000	0
399	15594041	Female	49	36000	1

400 rows × 5 columns

In [4]:

df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 400 entries, 0 to 399
Data columns (total 5 columns):
#   Column                Non-Null Count  Dtype
---  -
0   User ID                400 non-null   int64
1   Gender                 400 non-null   object
2   Age                    400 non-null   int64
3   EstimatedSalary        400 non-null   int64
4   Purchased              400 non-null   int64
dtypes: int64(4), object(1)
memory usage: 15.8+ KB
```

```
In [5]: df.isnull().sum()
```

```
# no null values;
```

```
Out[5]: User ID          0  
Gender          0  
Age            0  
EstimatedSalary 0  
Purchased       0  
dtype: int64
```

```
In [6]: df.shape
```

```
Out[6]: (400, 5)
```

```
In [9]: df.describe()
```

```
Out[9]:
```

	User ID	Age	EstimatedSalary	Purchased
count	4.000000e+02	400.000000	400.000000	400.000000
mean	1.569154e+07	37.655000	69742.500000	0.357500
std	7.165832e+04	10.482877	34096.960282	0.479864
min	1.556669e+07	18.000000	15000.000000	0.000000
25%	1.562676e+07	29.750000	43000.000000	0.000000
50%	1.569434e+07	37.000000	70000.000000	0.000000
75%	1.575036e+07	46.000000	88000.000000	1.000000
max	1.581524e+07	60.000000	150000.000000	1.000000

```
In [12]: X = df.iloc[:, [2, 3]].values  
# X = df[['EstimatedSalary', 'Age']]  
y = df.iloc[:, 4].values  
# y = df['Purchased']
```

Test-size = 0.2

```
In [63]: # Splitting the dataset into the Training set and Test set

from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2, random_state = 42)
```

```
In [64]: # Feature Scaling
from sklearn.preprocessing import StandardScaler
sc = StandardScaler()
X_train = sc.fit_transform(X_train)
X_test = sc.transform(X_test)
```

```
In [65]: from sklearn.linear_model import LogisticRegression
lr = LogisticRegression(random_state = 42)
lr.fit(X_train, y_train)
```

```
Out[65]: LogisticRegression
LogisticRegression(random_state=42)
```

```
In [66]: y_pred = lr.predict(X_test)
```

```
In [67]: y_pred
```

```
Out[67]: array([0, 1, 0, 1, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 1, 0, 1, 0, 0,
                0, 1, 0, 0, 1, 0, 1, 0, 0, 0, 1, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0,
                0, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0,
                1, 1, 0, 0, 1, 0, 0, 0, 0, 0, 1, 1, 0, 0])
```

```
In [68]: y_test
```

```
Out[68]: array([0, 1, 0, 1, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0,
                1, 1, 0, 1, 0, 0, 1, 0, 1, 0, 1, 0, 1, 0, 0, 0, 0, 0, 1, 0, 0, 1,
                0, 1, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 1, 0, 0, 1,
                1, 1, 0, 0, 1, 0, 0, 0, 1, 0, 1, 1, 0, 1])
```


```
In [69]: from sklearn.metrics import confusion_matrix  
cm = confusion_matrix(y_test, y_pred)  
print(cm)
```

```
# TP FP  
# FN TN
```

```
[[50  2]  
 [ 9 19]]
```

```
In [70]: sns.heatmap(cm ,annot=True)
```

```
Out[70]: <AxesSubplot: >
```



```
In [71]: from sklearn.metrics import classification_report
print(classification_report(y_test,y_pred))
```

	precision	recall	f1-score	support
0	0.85	0.96	0.90	52
1	0.90	0.68	0.78	28
accuracy			0.86	80
macro avg	0.88	0.82	0.84	80
weighted avg	0.87	0.86	0.86	80

```
In [72]: from sklearn.metrics import accuracy_score
accuracy = accuracy_score(y_test, y_pred)
accuracy
```

Out[72]: 0.8625

```
In [73]: error_rate = 1-accuracy
error_rate
```

Out[73]: 0.13749999999999996

Test-size = 0.25

```
In [39]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25, random_state=25)
```

```
In [40]: X_train = sc.fit_transform(X_train)
X_test = sc.transform(X_test)
```

```
In [41]: lr = LogisticRegression(random_state=25)
         lr.fit(X_train, y_train)
```

```
Out[41]: LogisticRegression
         LogisticRegression(random_state=25)
```

```
In [42]: y_pred = lr.predict(X_test)
```

```
In [43]: cm = confusion_matrix(y_test, y_pred)
```

```
In [44]: print(cm)
```

```
# TP FP
# FN TN
```

```
[[58  8]
 [13 21]]
```

```
In [45]: print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.82	0.88	0.85	66
1	0.72	0.62	0.67	34
accuracy			0.79	100
macro avg	0.77	0.75	0.76	100
weighted avg	0.79	0.79	0.79	100

```
In [ ]: from sklearn.metrics import accuracy_score
         accuracy = accuracy_score(y_test, y_pred)
         accuracy
```

Test-size = 0.3

```
In [52]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=25)
```

```
In [53]: X_train = sc.fit_transform(X_train)
X_test = sc.transform(X_test)
```

```
In [54]: lr = LogisticRegression(random_state=30)
lr.fit(X_train, y_train)
```

```
Out[54]: ▼      LogisticRegression
LogisticRegression(random_state=30)
```

```
In [55]: y_pred = lr.predict(X_test)
```

```
In [56]: cm = confusion_matrix(y_test, y_pred)
```

```
In [57]: print(cm)
```

```
# TP FP
# FN TN
```

```
[[70  7]
 [16 27]]
```

```
In [58]: print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.81	0.91	0.86	77
1	0.79	0.63	0.70	43
accuracy			0.81	120
macro avg	0.80	0.77	0.78	120
weighted avg	0.81	0.81	0.80	120


```
In [59]: from sklearn.metrics import accuracy_score  
accuracy = accuracy_score(y_test, y_pred)  
accuracy
```

Out[59]: 0.8083333333333333

```
In [61]: error_rate = 1-accuracy  
error_rate
```

Out[61]: 0.19166666666666665

```
In [ ]:
```